

Learning the topology of graphs

Convex Group-HKUST
2018-08-17

Contents

1	Problem Statement	1
2	Usage of the package	1
3	Explanation of the algorithms	4
3.1	learnGraphTopology: Learning the topology of graph	5
	References	5

1 Problem Statement

The problem of learning a K-component graph may be expressed mathematically as

$$\underset{\mathbf{w}, \mathbf{\Lambda}, \mathbf{U}}{\text{minimize}} -\log \det(\mathbf{\Lambda}) + \text{tr}(\mathbf{K}\mathcal{L}\mathbf{w}) + \frac{\beta}{2} \|\mathcal{L}\mathbf{w} - \mathbf{U}\mathbf{\Lambda}\mathbf{U}^T\|_F^2 \quad (1)$$

$$\text{subject to } \mathbf{w} \geq 0, \mathbf{\Lambda} \in \mathcal{S}_{\mathbf{\Lambda}}, \mathbf{U}^T \mathbf{U} = \mathbf{I}, \quad (2)$$

in which $\mathcal{S}_{\mathbf{\Lambda}}$ is the space of eigenvalues of matrices $\mathbf{\Theta}$ which are positive semi-definite, symmetric, and whose sum of the elements of any row or column is equal to zero.

We use a block coordinate descent to optimize each variable while holding the others fixed. Therefore, problem (??) can be divided into the following problems:

2 Usage of the package

We illustrate the usage of the package with simulated data, as follows:

```
library(spectralGraphTopology)
set.seed(123)

# Number of samples
T <- 10000
# Vector to generate the Laplacian matrix of the graph
w <- runif(10)
# Laplacian matrix
Theta <- L(w)
# Sample data from a Multivariate Gaussian
N <- ncol(Theta)
Y <- MASS::mvrnorm(T, rep(0, N), MASS::ginv(Theta))
# Number of components of the graph
K <- 1
# Learn the Laplacian matrix
res <- learnGraphTopology(Y, K, beta = 10)
```

Let's visually inspect the true Laplacian and the estimated one:

```

Theta
#>           [,1]      [,2]      [,3]      [,4]      [,5]
#> [1,]  2.3678770 -0.2875775 -0.7883051 -0.4089769 -0.8830174
#> [2,] -0.2875775  1.8017068 -0.9404673 -0.0455565 -0.5281055
#> [3,] -0.7883051 -0.9404673  3.1726265 -0.8924190 -0.5514350
#> [4,] -0.4089769 -0.0455565 -0.8924190  1.8035672 -0.4566147
#> [5,] -0.8830174 -0.5281055 -0.5514350 -0.4566147  2.4191726
res$Theta
#>           [,1]      [,2]      [,3]      [,4]      [,5]
#> [1,]  2.3384299 -0.29372890 -0.7864202 -0.38899191 -0.8692889
#> [2,] -0.2937289  1.76412935 -0.9297244 -0.05376403 -0.4869120
#> [3,] -0.7864202 -0.92972444  3.0917157 -0.81448289 -0.5610881
#> [4,] -0.3889919 -0.05376403 -0.8144829  1.74104811 -0.4838093
#> [5,] -0.8692889 -0.48691199 -0.5610881 -0.48380928  2.4010983

```

We can evaluate the performance of the learning process in a more objective manner by computing the relative error between the true Laplacian matrix and the estimated one, which can be done as follows::

```

RE <- norm(Theta - res$Theta, type="F") / max(1., norm(Theta, type="F"))
RE
#> [1] 0.02965113


```

Let's also look at the trend of the objective function per iteration:

```

k <- length(res$fun)
plot(c(1:k), res$fun, pch=19, cex=.5, col = scales::alpha("black", .25),
     xlab = "Iteration number", ylab = "Objective function")

```



For $K > 1$, we can generate the Laplacian as a block diagonal matrix, as follows

```

library(spectralGraphTopology)
w1 <- runif(3)
w2 <- runif(3)
Theta1 <- L(w1)
Theta2 <- L(w2)
N1 <- ncol(Theta1)

```

```

N2 <- ncol(Theta2)
Theta <- rbind(cbind(Theta1, matrix(0, N1, N2)),
               cbind(matrix(0, N2, N1), Theta2))
Y <- MASS::mvrnorm(T, rep(0, N1 + N2), MASS::ginv(Theta))
K <- 2
beta <- 5
res <- learnGraphTopology(Y, K, beta = beta, ftol = 1e-3)
RE <- norm(Theta - res$Theta, type="F") / max(1., norm(Theta, type="F"))
RE
#> [1] 0.8834062

```

```

Theta
#>           [,1]      [,2]      [,3]      [,4]      [,5]      [,6]
#> [1,]  0.96494812 -0.03715792 -0.9277902  0.0000000  0.0000000  0.0000000
#> [2,] -0.03715792  0.48514753 -0.4479896  0.0000000  0.0000000  0.0000000
#> [3,] -0.92779020 -0.44798960  1.3757798  0.0000000  0.0000000  0.0000000
#> [4,]  0.00000000  0.00000000  0.0000000  1.2716541 -0.9159023 -0.3557518
#> [5,]  0.00000000  0.00000000  0.0000000 -0.9159023  1.8312150 -0.9153127
#> [6,]  0.00000000  0.00000000  0.0000000 -0.3557518 -0.9153127  1.2710645
res$Theta
#>           [,1]      [,2]      [,3]      [,4] [,5]      [,6]
#> [1,]  1.3891578  0.0000000 -0.3605830 -0.6061550  0 -0.4224198
#> [2,]  0.0000000  0.7412041  0.0000000 -0.3010095  0 -0.4401947
#> [3,] -0.3605830  0.0000000  1.6932700 -0.5816278  0 -0.7510592
#> [4,] -0.6061550 -0.3010095 -0.5816278  1.6025191  0 -0.1137268
#> [5,]  0.0000000  0.0000000  0.0000000  0.0000000  0  0.0000000
#> [6,] -0.4224198 -0.4401947 -0.7510592 -0.1137268  0  1.7274005

```

As we can observe, the matrices' structure do not quite match.
I suspect they might be similar matrices. Let's check some properties:

```

eigen(Theta, only.values = TRUE)
#> $values
#> [1]  2.746823e+00  2.185018e+00  1.627111e+00  6.408571e-01 -3.252607e-18
#> [6] -1.948853e-17
#>
#> $vectors
#> NULL
eigen(res$Theta, only.values = TRUE)
#> $values
#> [1]  2.713545e+00  1.965910e+00  1.661742e+00  8.123542e-01  0.000000e+00
#> [6] -2.220446e-16
#>
#> $vectors
#> NULL

```

```

sum(diag(Theta))
#> [1] 7.199809
sum(diag(res$Theta))
#> [1] 7.153552

```

```

det(Theta)
#> [1] 0
det(res$Theta)
#> [1] 0


N <- N1 + N2
evd_true <- eigen(Theta)
vec_true <- evd_true$vectors[, N:1]
vec_true
#>           [,1]      [,2]      [,3]      [,4]      [,5]      [,6]
#> [1,]  0.0000000 -0.5773503  0.5643095  0.0000000000  0.5901029  0.0000000
#> [2,]  0.0000000 -0.5773503 -0.7931989  0.0000000000  0.1936549  0.0000000
#> [3,]  0.0000000 -0.5773503  0.2288893  0.0000000000 -0.7837578  0.0000000
#> [4,] -0.5773503  0.0000000  0.0000000 -0.7069205454  0.0000000  0.4085707
#> [5,] -0.5773503  0.0000000  0.0000000 -0.0003723246  0.0000000 -0.8164965
#> [6,] -0.5773503  0.0000000  0.0000000  0.7072928700  0.0000000  0.4079258
vals_true <- evd_true$values[N:1]
vals_true
#> [1] 8.881784e-16 1.332268e-15 6.408571e-01 1.627111e+00 2.185018e+00
#> [6] 2.746823e+00
res$lambda
#> [1] 0.000000 0.000000 1.010313 1.774452 2.062863 2.785350
objFunction(Theta, vec_true, vals_true, res$Km, beta, N, K)
#> [1] 2.185048
objFunction(res$Theta, vec_true, vals_true, res$Km, beta, N, K)
#> [1] 31.74898

```

```

k <- length(res$fun)
plot(c(1:k), res$fun, pch=19, cex=.6, xlab = "Iteration number",
     ylab = "Objective function")

```



3 Explanation of the algorithms

In this section we describe in detail the algorithms designed to solve the graph topology learning problem.

3.1 learnGraphTopology: Learning the topology of graph

The goal of `learnGraphTopology()` is to estimate the Laplacian matrix generated by the weight vector of a graph, \mathbf{w} . The algorithm for the function `learnGraphTopology` is stated as follows:

Data: \mathbf{Y} (data matrix), K ($\#\{\text{components}\}$), β (regularization term), $\mathbf{w}_0, \boldsymbol{\lambda}_0, \mathbf{U}_0$ (initial parameter estimates), α_1, α_2 (lower and upper bound on the eigenvalues of the Laplacian matrix), ρ (how much to increase beta per iteration)

Result: $\boldsymbol{\Theta}$ (Laplacian matrix)

$N \leftarrow \text{ncol}(\mathbf{Y})$

while *objective function do not converged or max $\#\{\text{iterations}\}$ not reached* **do**

$k \leftarrow 0$

while *parameters do not converged or max $\#\{\text{iterations}\}$ not reached* **do**

$\mathbf{w}^{(k+1)} \leftarrow \text{w_update}(\mathbf{w}^{(k)}, \mathbf{U}^{(k)}, \boldsymbol{\lambda}^{(k)}, \beta, N, \mathbf{K})$

$\mathbf{U}^{(k+1)} \leftarrow \text{U_update}(\mathbf{w}^{(k+1)}, N)$

$\boldsymbol{\lambda}^{(k+1)} \leftarrow \text{lambda_update}(\mathbf{w}^{(k+1)}, \mathbf{U}^{(k+1)}, \alpha_1, \alpha_2, \beta, N, K)$

$k \leftarrow k + 1$

end

$\beta \leftarrow \beta(\rho + 1)$

end

return $\mathcal{L}(\mathbf{w}^{(k+1)})$

Function `w_update(w, U, λ, β, N, K):`

$\nabla_{\mathbf{w}} f \leftarrow \mathcal{L}^* \left(\mathcal{L}(\mathbf{w}) - \mathbf{U} \text{diag}(\boldsymbol{\lambda}) \mathbf{U}^T + \frac{\mathbf{K}}{\beta} \right)$

return $\max \left(0, \mathbf{w} - \frac{\nabla_{\mathbf{w}} f}{2N} \right)$

Function `U_update(w, N):`

return `eigen(L(w))$vectors[, N : 1]`

Function `lambda_update(w, U, α1, α2, β, N, K):`

$\mathbf{d} \leftarrow \text{diag}(\mathbf{U}^T \mathcal{L}(\mathbf{w}) \mathbf{U})$

$\boldsymbol{\lambda} \leftarrow \frac{1}{2} \left(\mathbf{d} + \sqrt{\mathbf{d} \odot \mathbf{d} + \frac{4}{\beta}} \right)$

if $\boldsymbol{\lambda}$ *has its elements in increasing order* **then**

return $\boldsymbol{\lambda}$

else

 set to α_1 the elements of $\boldsymbol{\lambda}$ whose values are less than α_1

 set to α_2 the elements of $\boldsymbol{\lambda}$ whose values are greater than α_2

end

if $\boldsymbol{\lambda}$ *has its elements in increasing order* **then**

return $\boldsymbol{\lambda}$

else

raise `Exception("eigenvalues are not in increasing order")`

end

References